

Hello and thank you for agreeing to help review our study materials!

Our project examines how popular LLMs respond when users solicit information about their legal situations. We believe that frontier AI labs have quietly lifted guardrails that discouraged chatbots from dispensing legal advice. Such a change would be troubling, since these systems are not held to the standards of competence and ethical behavior enforced by the legal profession. To determine whether LLMs are now performing tasks that have historically required bar admission, we operationalize the boundary around legal practice using Unauthorized Practice of Law (UPL) case law.

A law student on our team has collected and briefed a set of 80 UPL cases, summarizing the facts, legal reasoning, outcomes, and harms to the people who had received improper advice. We draw cases from five high-stakes domains where outcomes are difficult to reverse.

Your expertise is vital to our project – are we adequately representing your domain?

As a subject matter expert, we would appreciate your help with the following:

Read the “Prompt Components” and “Rubric” for each case in your domain & tell us if you agree with the application of the rubric to chat transcripts

- Feel free provide feedback on the prompts or rubrics to better capture judicial reasoning or your experience in this domain
 - Any changes you suggest to the prompts will change our inputs into the chatbot models.
- We include five sample transcripts for each case (randomly selected from our set of chatbot models) so that you can see some variation in model behavior.
 - Have we adequately characterized whether the chatbot in each transcript has engaged in UPL?

Review the quality of the legal service and advice in the chat transcripts

- What is useful in these transcripts?
- What is troubling?
- We are interested in both:
 - Your opinion about the **quality of the legal information/advice** provided
 - Does the model get the law right in this domain?
 - Your opinion about the **quality of the “legal service”**
 - What aspects of the transcripts could be improved in this situation?
 - E.g., we have observed harmful patterns in some chatbot transcripts:
 - Insufficient interviewing: failure to surface relevant information prior to giving legal advice
 - For instance, failing to ask about past deportations before recommending the user to file a I-130
 - Lack of jurisdiction / Incorrect inferred jurisdiction

- Failures of analysis
 - Failures to cite correct case law or other authorities
- What aspects of the transcripts could be useful to people in this situation?
 - E.g., we have observed some useful recommendations
 - Directions to consult with experts
 - Resources / links to connect with aid organizations
 - Eliciting additional information
- Feel free to include any other observations -- what should we know about that we haven't asked for yet?
 - Please note that we will be anonymizing personal details from cases during our final data collection phase and in future publications, but include them here to support traceability of prompt information to case information.

Please note that we will be anonymizing personal details from cases during our final data collection phase and in future publications, but include them here to support traceability of prompt information to case information.

If you are curious about how we construct our benchmark or rubrics, we outline our approach below.

Our approach to composing the benchmark:

- We collected 80 cases where UPL was found.
- We created one set of queries per case.
 - These three queries are issued sequentially to establish a multi-turn conversation.
- We wanted to emulate the person who had received the improper legal advice, as our project imagines them as now asking an LLM for advice.
- First query components:
 - Facts: We first composed a "facts blurb." We compiled information about the person who had received the improper legal advice and their situation. We then wrote a description of their initial legal troubles, as told from their POV.
 - Jurisdiction. We then wrote a sentence about their location
 - (This will be blank in half of the tested queries)
 - Prompt/Task: We requested that the LLM respond to some query or perform some task. We phrase these requests to specifically solicit advice or work that the court called UPL.
 - (e.g., the court says drafting a pleading for someone else is UPL, we request the LLM to draft a pleading)
 - We concatenate facts + jurisdiction + prompt as the first message to the LLM
- Second query components:
 - Follow-up 1: We make another request of the LLM, again soliciting advice or

work that the court has marked UPL.

- Third query components:
 - Follow-up 2: We make another request of the LLM, again soliciting advice or work that the court has marked UPL.

Our approach to composing the rubrics:

- The rubrics serve as a guide to the annotators. They help decide whether a particular response would constitute UPL according to the court's logic in the case.
 - The annotators select a binary answer (Yes/No) to indicate whether a particular response gives legal advice
 - We create a rubric for each component of our query, such that there are three rubrics per case
 - We took a first pass on composing the rubric based on the court reasoning.
 - (e.g., the court says drafting a pleading for someone else is UPL, so we consider it UPL if the response includes language the model recommends including in a pleading)
-